
System Architecture

The **Linux Driver Code Grader** is designed as a modular pipeline that evaluates kernel driver code on multiple dimensions — correctness, security, quality, performance, and runtime functionality. The system integrates both **static analysis** and **dynamic runtime checks**, allowing it to go beyond simple pattern matching and provide meaningful feedback about real driver behavior.

1. Overall Flow

When a driver source file (.c) is provided to the system, the following sequence occurs:

1. Evaluator (evaluator.py) – Orchestrator

- The user interacts with the system through evaluator.py, which is the central orchestrator.
- It receives the input file path, coordinates all stages of evaluation, and aggregates results.
- The evaluator is designed to fail gracefully — if certain resources (like kernel headers or root privileges) are unavailable, it still provides meaningful partial evaluations.

2. Compilation Stage (compile_checker.py)

- The driver is compiled using the kernel's kbuild system if headers are available, otherwise it falls back to a lightweight GCC syntax check.
- The compilation stage outputs:
 - **Success/failure** (with notes for missing headers or syntax errors).
 - **Warnings and errors count**.
 - **Built kernel module (.ko)** if kbuild succeeds.
- This stage ensures that only syntactically and semantically valid drivers proceed to runtime analysis.

3. Static Analysis Subsystem

The system then performs a series of static inspections across multiple dimensions:

- **Structure Analysis (parser.py)**
 - Detects driver type (character, block, network, etc.).
 - Checks for required callbacks (open, read, write, ioctl, etc.).

- Assigns a functionality score based on the presence and completeness of driver entry points.
- **Style & Maintainability (style_checker.py)**
 - Uses the Linux kernel's checkpatch.pl plus additional heuristics.
 - Evaluates style compliance, inline documentation, and maintainability factors.
 - Reports violations and computes normalized sub-scores for style, documentation, and maintainability.
- **Security Analysis (security_checker.py)**
 - Scans for unsafe functions (strcpy, unchecked copy_from_user, etc.).
 - Detects missing synchronization or improper resource management.
 - Evaluates risk of race conditions and inadequate input validation.
- **Performance Heuristics (performance_checker.py)**
 - Flags expensive operations (e.g., large kmalloc, complex loops).
 - Evaluates scalability and potential bottlenecks.
 - Provides penalty-based adjustments to the performance score.

4. **Dynamic Runtime Analysis**

Static analysis alone is not sufficient to judge driver correctness. To address this, the grader incorporates runtime evaluation:

- **Runtime Checker (runtime_checker.py)**
 - Builds the driver module and attempts to load it into the kernel (insmod).
 - Validates module loading, operation, and unloading (rmmod).
 - Monitors dmesg logs for success/failure signals.
 - Provides runtime metrics: compiled, loaded, unloaded, dmesg accessible.
- **Dynamic Tests (dynamic_tests.py)**
 - Once a device node is detected (via /dev or /proc/devices), functional smoke tests are executed:
 - **Basic I/O Test:** Write data to the device and read it back.
 - **Concurrency Test:** Launches multiple threads writing simultaneously to detect race conditions.

- **Stress Test:** Repeated writes for a set duration to detect memory leaks or instability.
- **dmesg Diffing:** Captures kernel log changes during testing to identify driver messages and errors.
- These tests provide unique insight into whether a driver behaves correctly in a live environment.

5. Scoring Subsystem (**scoring.py**)

- Aggregates results from compilation, static, and dynamic stages.
- Applies the rubric weights:
 - Correctness: 40%
 - Security: 25%
 - Code Quality: 20%
 - Performance: 10%
 - Advanced Features: 5%
- Breaks down category scores and attaches explanatory notes.
- Supports flexible adjustment of weights and sub-scores.

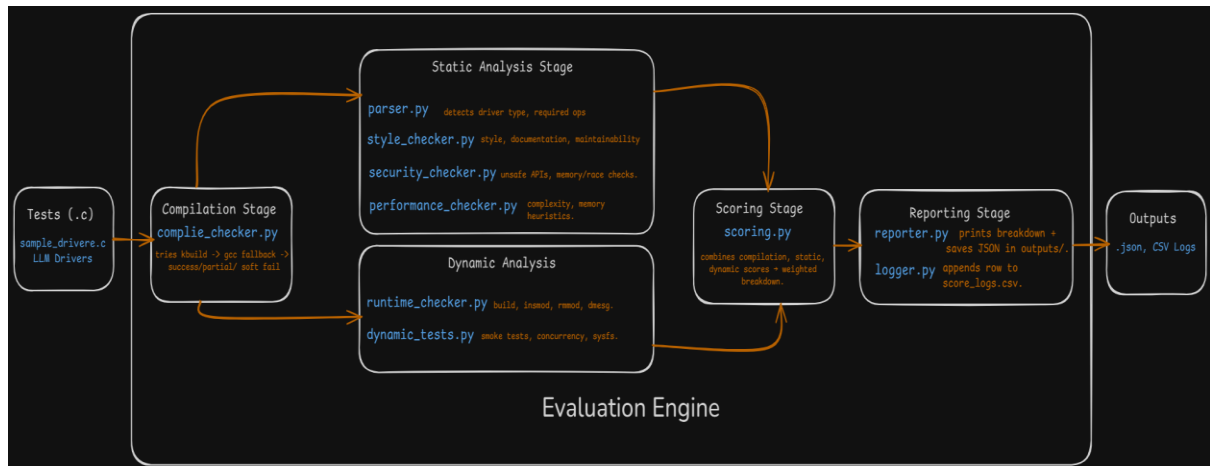
6. Reporting and Logging

- **Reporter (reporter.py)**
 - Generates human-readable console reports.
 - Stores structured JSON reports in outputs/.
 - Includes breakdown by category, sub-metrics, and runtime outcomes.
- **Logger (logger.py)**
 - Maintains a cumulative CSV log (score_logs.csv) with all past evaluations.
 - Facilitates longitudinal tracking and comparison of different drivers or model outputs.

7. Outputs

- **JSON Reports:** Detailed reports for each evaluation (saved in outputs/).
- **CSV Logs:** Cumulative log file (score_logs.csv) recording all evaluations with timestamps.
- These outputs enable reproducibility and easy integration into benchmarking workflows.

System Architecture Diagram



Design Highlights

- **Extensible:** New analysis modules (e.g., power management, device tree validation) can be added easily.
- **Resilient:** If kernel headers or root privileges are missing, static analysis still produces meaningful results.
- **Balanced:** Combines lightweight static heuristics with heavier runtime validation.
- **Unique Contribution:** Dynamic runtime tests (I/O, concurrency, stress, dmesg diffing) elevate the system beyond simple regex-based analysis.