

---

# User Guide for Linux-Driver-Code-Grader

This document explains how to install, run, and analyze evaluations of Linux driver source files using the **Linux-Driver-Code-Grader** system.

---

## 1. Prerequisites

The evaluator requires a Linux environment with kernel headers. Supported environments:

- **Ubuntu 22.04+** (recommended)
- **GitHub Codespaces** (tested)
- **Other Linux distributions** with minimal adjustments

### Required Packages

- Python 3.9+
  - build-essential (gcc, make)
  - Linux kernel headers (/lib/modules/\$(uname -r)/build)
  - Static analysis tools:
    - sparse
    - clang-format, clang-tidy
    - cppcheck
  - gdb
  - kmod (for insmod, rmmod, lsmod)
- 

## 2. Installation

### Clone the repository

```
git clone https://github.com/VedantNipane/Linux-Driver-Code-Grader.git
```

```
cd Linux-Driver-Code-Grader
```

### Setup environment

The repo provides a helper script:

```
chmod +x setup.sh
```

```
./setup.sh
```

This installs all necessary dependencies (requires sudo).

---

### 3. Running Evaluations

#### Evaluate a single driver

```
python3 Evaluator/evaluator.py Tests/sample_driver.c
```

This will:

1. Attempt compilation (kbuild first, fallback to gcc).
  2. Run static analysis (style, parser, security, performance).
  3. Run dynamic tests (load/unload, dmesg checks, smoke tests).
  4. Generate a JSON report under outputs/ and append results to score\_logs.csv.
- 

#### Run evaluation on all test drivers

Use the provided Makefile:

```
make run
```

This executes evaluator.py against every file in Tests/ and collects outputs.

---

### 4. Understanding Outputs

#### Console Report

Each run prints a structured report:

```
=== Evaluation Report ===
```

```
File: Tests/sample_driver.c
```

```
Compilation: Success (method=kbuild)
```

```
Warnings: 1 Errors: 0
```

```
--- Score Breakdown ---
```

```
Correctness: 38.0/40
```

- Compilation: success (method=kbuild)
- Functionality score: 1.00 -> 10.00/10
- Runtime score: 8.0/10 (compiled, loaded, unloaded)

Security: 23.1/25

- sub\_scores: {memory\_safety: 0.7, resource\_mgmt: 1.0, ...}

- issues: ['unsafe\_function:strcpy']

Code Quality: 13.2/20

Performance: 8.0/10

Advanced: 0.0/5

Overall Score: 82.3/100

Report saved to: outputs/sample\_driver\_results.json

---

## JSON Report

Stored under outputs/filename\_results.json.

Contains structured data:

```
{  
  "file": "Tests/sample_driver.c",  
  "overall_score": 82.3,  
  "breakdown": {  
    "Correctness": {"awarded": 38.0, "max": 40, "details": [...]},  
    "Security": {"awarded": 23.1, "max": 25, "details": [...]}  
  },  
  "compilation": {...},  
  "runtime": {...},  
  "security": {...},  
  "style": {...},  
  "performance": {...}  
}
```

This format is machine-readable for downstream analysis.

---

## CSV Log

All evaluations append results to `score_logs.csv` (repo root).

Example entry:

```
timestamp,file,overall_score,Correctness,Security,Code Quality,Performance,Advanced
2025-09-18 17:42:03,sample_driver.c,77.5/100,34.0/40,18.8/25,16.7/20,8.0/10,0.0/5
```

Useful for tracking trends across multiple runs.

---

## 5. Technical Details

### Compilation Modes

- **kbuild (preferred):** Uses Linux kernel's native build system. Produces `.ko` modules if successful.
- **gcc fallback:** Syntax-only check if headers are missing or `kbuild` fails. May yield a "soft pass".

### Dynamic Tests

- Module insertion/removal via `insmod / rmmod`
- Kernel log check via `dmesg`
- Smoke test for device node presence in `/dev` or `/proc/devices`
- Output integrated into **Correctness score**

### Scoring

- Correctness capped at 40 (compilation + structure + runtime)
  - Security normalized from multiple sub-checks
  - Performance starts at full score, deductions applied
  - Advanced features (`devm`, device tree, pm hooks, debugfs) give bonus points
- 

## 6. Limitations

- Requires `root/sudo` for runtime tests (`insmod`, `rmmod`).
  - Kernel headers must match the running kernel version.
  - Smoke tests do not validate driver functionality beyond load/unload and node creation.
  - Performance heuristics are static (no stress testing).
- 

## 7. Next Steps

- Extend dynamic tests (I/O validation, stress tests).
  - Broader driver type support (block, network).
  - Integration with CI pipelines for automated evaluation.
-